

Individual Project: Build a Personal Research Portal (PRP)

Three phases with progressive deliverables: Prompting → RAG → Research Portal Product

Work mode	Individual
Grade weight	Phase 1: 10% • Phase 2: 15% • Phase 3: 20% (Total: 45% of course grade)
Primary goal	A research-grade portal that answers questions with evidence, citations, and exportable research artifacts
Stack	Your choice (Python + Streamlit/Gradio recommended). Free-tier and open-weight tools allowed.
Submission	GitHub repo (code + data manifest) + reports + demo recording

1. Overview (what you are building)

By the end of Phase 3 you will deliver a Personal Research Portal (PRP): a lightweight product that helps you move from a research question to a grounded synthesis. It ingests a domain corpus, retrieves evidence, produces citation-backed answers and research artifacts, and logs evaluation results.

Minimum PRP capabilities (MVP)

- Ingest a domain corpus (PDFs/URLs/notes) with metadata (author, year, source type, link/DOI).
- Retrieve evidence (semantic retrieval at minimum).
- Generate answers where every major claim is backed by an inline citation to a source/chunk ID.
- Produce at least one research artifact: evidence table, annotated bibliography, or synthesis memo.

- Save research threads (query, retrieved chunks, answer) and export outputs (Markdown/CSV/PDF).
- Run an evaluation set and report groundedness/faithfulness plus at least one additional metric.

What counts as a citation in this project

- You may cite using (source_id) or (source_id, chunk_id) as long as it uniquely maps to text you ingested.
- Your repo must include a data manifest that makes every citation resolvable to: source title + link/DOI + local file or snapshot.
- Example citation formats: (Smith2023, chunk_07) or (Smith2023) if the source_id uniquely identifies one ingested document.
- If the corpus does not support a claim, your system must say so rather than guessing.

2. How the three phases fit together

- Phase 1 produces your research framing, prompt kit, and an evaluation rubric you will reuse.
- Phase 2 uses that same framing to build research-grade retrieval and grounding over your corpus.
- Phase 3 wraps Phase 2 into a usable portal and generates exportable research artifacts.

Design intent: each phase should feel like a direct upgrade of the previous one, not a reset.

3. Deliverables at a glance

Phase	Weight	Purpose	Core deliverables	MVP scope (minimum)	Stretch goals (examples)
Phase 1 – Prompt the Research Domain	10%	Design prompts and evaluate model behavior on research tasks.	Framing brief; prompt kit; evaluation sheet; analysis memo.	2 tasks • 2 models • 2 test cases per task • 2 prompt variants per task (16 total runs).	3rd model; adversarial tests; automated checks; better uncertainty calibration.

Phase 2 – Ground the Domain (Research- Grade RAG)	15%	Build baseline RAG over your corpus; add production patterns; evaluate.	Ingestion pipeline; baseline RAG; enhancements; evaluation report + logs.	15 sources • 20 queries • 1 enhancement	Hybrid retrieval; reranking; section-aware chunking; structured citations; RAGAs/ARES.
Phase 3 – Build the Personal Research Portal Product	20%	Turn your RAG into a usable portal that produces research artifacts.	Working app + demo; final report; generated artifacts; evaluation view.	Search + citations + threads + export + at least 1 artifact type.	Agentic research loop; knowledge graph; gap finding; BibTeX export; improved UX.

4. Phase 1 – Prompt the Research Domain (10%)

Goal

Pick a research domain, define a main research question, and test how different models handle core research tasks. You will produce a prompt kit you will reuse in Phases 2 and 3.

You must choose

- A domain and one main research question.
- 4–6 sub-questions that break your main question into retrievable parts.
- Two tasks from the task menu (each with a required output format).
- Two accessible models (three allowed for stretch).

Task menu (choose 2) – with concrete output formats

- Paper triage: output a 5-field summary (Contribution, Method, Data, Findings, Limitations).
- Claim-evidence extraction: output 5 rows with Claim | Direct quote/snippet | Citation (source_id, chunk_id).
- Cross-source synthesis: output a table of Agreement | Disagreement | What evidence supports each side.

- Research gap spotting: output 3 gaps, each with 'What is missing' + 'What evidence would resolve it'.
- Citation formatting: output a paragraph with inline citations plus a short reference list from your manifest.

Definitions (what we mean in Phase 1)

- Task: a repeatable research operation with a required output format (e.g., paper triage or claim-evidence extraction).
- Test case: one concrete input for a task (e.g., Paper A PDF; or Paper B, Methods section). You will run every test case through both prompts and both models.
- Prompt variant: two versions per task - Prompt A (baseline) and Prompt B (improved with structure/guardrails).

Worked example (illustrative - you will use your own domain)

- Domain + question: "Trustworthy RAG evaluation methods" - Main question: "How do different faithfulness metrics fail, and how can we combine them?"
- Tasks (pick 2): (1) Paper triage; (2) Claim-evidence extraction.
- Test cases (2 per task):
 - Paper triage test cases: Paper A (survey) and Paper B (empirical study).
 - Claim-evidence extraction test cases: Paper C (benchmark) and Paper D (critique).
- Prompts (2 per task): Prompt A = minimal instruction; Prompt B = structured output + "cite chunk_id" + "say unknown when missing".
- Models (2): any two you can access. Run all combinations (16 runs for MVP) and score outputs for groundedness and citation correctness.

MVP checklist (do these in order)

- Write a 1-page framing brief: domain, main question, sub-questions, scope (what you include/exclude).
- Create 2 test cases per chosen task (4 total test cases across Phase 1). A test case is one concrete input, such as a specific paper or section.
- Write two prompt variants per task: Prompt A (baseline) and Prompt B (improved structured prompt with guardrails).
- Run both models across all test cases and both prompt variants, and record outputs in the evaluation sheet (2 tasks x 2 cases x 2 prompts x 2 models = 16 runs).
- Score each output using the Phase 1 rubric and write a 1–2 page analysis of failure modes and fixes.

Phase 1 scoring rubric (1–4) – what the numbers mean

- 4: Correctly grounded and structured; citations are correct; uncertainty is stated when evidence is weak.
- 3: Mostly correct and structured; minor missing nuance OR minor citation/format issues.
- 2: Partially correct; key omissions OR weak grounding OR vague citations.
- 1: Not usable; hallucinated claims, fabricated citations, or fails the required structure.

Phase 1 deliverables (submit together)

- Framing brief (1–2 pages).
- Prompt kit (Markdown): prompts + why each constraint exists.
- Evaluation sheet (CSV/Google Sheet/Markdown table): scores + notes per run (for MVP, include 16 rows - one per run).
- Analysis memo (1–2 pages): patterns, failures, and Phase 2 design choices.

Common mistakes to avoid

- Choosing tasks that do not force citations (then you cannot assess grounding).
- Scoring without any written justification (graders cannot tell what you observed).
- Changing your question mid-phase (it breaks continuity into Phase 2).

5. Phase 2 – Ground the Domain (Research-Grade RAG) (15%)

Goal

Build a baseline RAG pipeline over your own corpus and upgrade it with production patterns. This phase prioritizes retrieval quality, traceability, and evaluation—not UI polish.

Corpus requirements (MVP)

- 15–30 sources minimum. At least 8 should be peer-reviewed papers, standards, or reputable technical reports.
- Every source must have metadata: source_id, title, authors, year, type, venue (if applicable), link/DOI, and a 1–2 sentence relevance note.
- Source acquisition can be manual (download PDFs) or automated (crawler/agentic selector). Either way, do not rely on live links alone: store the raw artifact (PDF/HTML snapshot/extracted text) under data/raw/ OR include a reproducible download script.
- If you use an automated crawler/agent to pick sources, include the code plus the run configuration (prompts/filters) and a short note describing how sources were selected (manual vs scripted vs agentic).

Baseline RAG MVP (do these in order)

- Ingest sources: parse and clean text; store text + metadata; create a data manifest (CSV/JSON).
- Chunk with a documented strategy (chunk size, overlap). Use section-aware chunking for papers if feasible.
- Embed and index in a vector store (FAISS acceptable).
- Retrieve top-k chunks per query and generate an answer that cites chunk/source IDs.
- Create an evaluation set of at least 20 queries: 10 direct, 5 synthesis/multi-hop, 5 ambiguity/edge cases.
- Example direct query: "What does <method/metric> measure, and what are its known failure modes?"
- Example synthesis query: "Compare Paper A vs Paper B on <criterion>; where do they agree and disagree?"
- Example edge-case query: "Does the corpus contain evidence for <claim>?" (If not found, your system must say so.)
- Evaluate groundedness/faithfulness plus one additional metric (answer relevance, context precision/recall, citation precision). Interpret results.

Required production patterns (must-have)

- Logging: store queries, retrieved chunks, model outputs, and prompt/version IDs.
- Reproducibility: pinned dependency versions; one-command run path (Makefile or a single script).
- Trust behavior: refuse to invent citations; explicitly flag missing or conflicting evidence.

Enhancements (implement at least 1 for MVP)

- Query rewriting or decomposition (question → sub-queries).
- Reranking (cross-encoder, LLM reranker, or heuristic reranker).
- Metadata filtering (year/author/type) and faceted retrieval.
- Hybrid retrieval (BM25 + vector).
- Confidence / evidence-strength scoring.
- Structured citations (inline citations + reference list from your manifest).

Phase 2 deliverables (submit together)

- Code + repo (baseline RAG) with README instructions.
- Data manifest (CSV/JSON) with metadata for every source.
- Evaluation report (3–5 pages): query set design, metrics, results, and what improved with enhancements.

- Run logs (machine-readable): at least 20 queries with retrieved chunks and outputs.

Phase 2 acceptance tests (what graders will check quickly)

- A single command produces: retrieval results, an answer with citations, and a saved log entry.
- Randomly sampled answers have citations that resolve to real source text in your corpus.
- Your evaluation report includes at least 3 representative failure cases with evidence.

6. Phase 3 – Build the Personal Research Portal Product (20%)

Goal

Turn your Phase 2 system into a usable portal that supports a research workflow: question → evidence → synthesis → export. Your portal should not just chat; it should produce research artifacts.

Product MVP requirements

- Interface: a working UI with at least: search, ask, show sources/citations, and history.
- Research threads: save query + retrieved evidence + answer (file-based OK).
- Artifact generator: produce at least one of: evidence table, annotated bibliography, or synthesis memo with references.
- Export: allow export of artifacts (Markdown/CSV/PDF).
- Evaluation: a page or script that runs your query set and summarizes metrics plus representative examples.
- Trust behavior: every answer includes citations; missing evidence is stated explicitly with a suggested next retrieval step.

Recommended artifact schemas (pick at least 1 for MVP)

- Evidence table: Claim | Evidence snippet | Citation (source_id, chunk_id) | Confidence | Notes.
- Annotated bibliography: 8–12 sources with 4 fields (claim, method, limitations, why it matters).
- Synthesis memo: 800–1200 words with inline citations and a reference list generated from your manifest.

Phase 3 deliverables (submit together)

- Working PRP app + instructions to run locally.

- Demo recording (3–6 minutes) showing retrieval, citations, artifact generation, and export.
- Final report (6–10 pages): architecture, design choices, evaluation, limitations, and next steps.
- Generated research artifact outputs included in the repo and referenced in the report.

Stretch goals (optional)

- Agentic research loop (plan → search → read → synthesize) with guardrails and logs.
- Knowledge graph view (entities/concepts linked to source passages).
- Automatic disagreement map (conflicts surfaced with citations).
- Gap finder: missing evidence + targeted next retrieval actions.
- Improved UX: filters by year/venue/type, reading list, tagging, saved collections.

7. Grading (points aligned to course weight)

Your project is graded on rigor, clarity, and evidence, not on paid tooling. Points below map directly to the phase weights.

Phase 1 (10 points = 10%)

- Framing quality and scope discipline (3): clear question, sub-questions, inclusions/exclusions.
- Prompt kit quality (4): structured prompts, guardrails, reusable prompt cards.
- Evaluation rigor and analysis (3): consistent scoring, failure tags, actionable takeaways.

Phase 2 (15 points = 15%)

- Corpus + manifest quality (4): metadata completeness, reproducibility, source credibility.
- Retrieval + grounding implementation (5): working RAG with correct citations and logs.
- Enhancement + measurable improvement (3): at least one enhancement with evidence it helped.
- Evaluation report quality (3): query set design, metrics, interpretation, failure cases.

Phase 3 (20 points = 20%)

- Portal MVP functionality (8): UI, threads, citations, export, reliability of core flow.

- Research artifacts (4): artifact schema correctness, usefulness, citation traceability.
- Evaluation + trust behaviors (4): run query set, summarize metrics, handles missing evidence.
- Engineering + communication (4): clean repo, run instructions, demo, clear final report.

8. Submission requirements

Repo expectations (minimum)

- A README that tells a grader how to run your system in 5 minutes or less.
- Pinned dependencies (requirements.txt or pyproject.toml).
- data/data_manifest.csv (or JSON) with the required metadata schema.
- Your repo must include data/raw/ with the downloaded PDFs/snapshots you used, OR a reproducible download script that regenerates data/raw/.
- A folder of outputs (artifacts/exports) generated by your system.
- Logs from evaluated runs (machine-readable).

Recommended repo structure

```

repo/
  README.md
  requirements.txt (or pyproject.toml)
  data/
    raw/ (downloaded PDFs, HTML snapshots, original notes)
    processed/ (parsed text, chunks, intermediate files)
    data_manifest.csv (or JSON)
  src/
    app/ (Phase 3 UI)
    ingest/ (parsers + chunking + optional download scripts)
    rag/ (retrieval + generation)
    eval/ (query sets, scripts, results)
    outputs/ (artifacts, exports)
    logs/ (runs, prompts, versions)
    report/ (Phase 1, Phase 2, Phase 3 writeups)

```

AI usage disclosure (required)

- Include a 1-page AI-usage log listing: tool name, what you used it for, and what you changed manually afterward.
- You are responsible for correctness, citations, and meeting requirements even if an AI tool suggested code or text.

Appendix A – Templates (copy/paste)

A1. Prompt card template

- Prompt name:
- Intent:
- Inputs (what you provide):
- Outputs (required structure):
- Constraints / guardrails (no fabricated citations; cite chunk IDs; etc.):
- When to use / when not to use:
- Failure modes to watch for:

A2. Evaluation row template (per answer)

- Task ID: | Test case ID: | Query ID:

Task ID:

- Prompt ID + model:
- Retrieved evidence IDs:
- Score 1 – Groundedness/faithfulness (1–4):
- Score 2 – Citation correctness (1–4):
- Score 3 – Usefulness (1–4):
- Notes + failure tag (missing evidence, wrong citation, overconfident, etc.):

A3. Source metadata schema (minimum fields)

source_id, title, authors, year, source_type, venue, url_or_doi, raw_path, processed_path, tags, relevance_note

raw_path should point to the file you store under data/raw/. processed_path is optional but recommended (e.g., parsed text under data/processed/).